

intermediate periods. As the time horizon is fixed, and saving implies a gross rate of return of exactly 1, the dimension of the state space satisfies  $|\mathcal{S}| = T \cdot (\bar{s} + 1)$ , while the action space satisfies  $|\mathcal{A}| = M + 1$ . Effectively, the set of *feasible* actions conditional on the state,  $s_t$ , satisfies  $|\mathcal{A}(s_t)| = s_t + 1 \leq \mathcal{A}$ .

With the purpose of learning the true value function  $V(s_t)$  for a given state  $s_t$ , we let our agent be an  $\epsilon$ -greedy Q-learner. In order to illustrate how the Q-learner learns the value function, it is instructive to define the *true value map* for the problem,  $Q(s_t, a_t) \mapsto \mathbb{R}$ , that maps a given state and action into a value as also discussed in the previous section. Hence, we may simply think of the relation between the value function and the value map as  $V(s_t) = \max_a Q(s_t, a_t)$ . That is, the value function implicitly selects the action  $a^*(s_t)$  in the set of feasible actions  $\mathcal{A}(s_t)$  that maximizes  $Q(s_t, a_t)$  holding  $s_t$  fixed.

In order to learn  $Q(s_t, a_t)$ , the agent is equipped with a matrix  $Q$  of size  $|\mathcal{S}| \cdot |\mathcal{A}|$  with all elements initialized to some values – in this case, we simply choose 0. Following an updating rule, the agent explores the environment and gradually updates the matrix  $Q$  in order to approximate  $Q(s_t, a_t)$ . As  $Q$  tends to  $Q(s_t, a_t)$ , the learner implicitly approximates the optimal actions for each state,  $s_t$ .

In the next subsection, we provide a simple example of a learning algorithm with an  $\epsilon$ -greedy Q-learner for the matrix  $Q$  that works well in discrete and bounded cases. As we move forward, we will introduce other methods for learning such representations, all of which either intend to approximate  $Q$ , intend to approximate the optimal action  $\mu^*(s_t)$ , or does both at the same time.


## 2.2 Updating the Q-matrix

With probability  $\epsilon$ , the agent chooses a random feasible action in the time- $t$  available action space  $\mathcal{A}(s_t) \in \mathcal{A}$ . This kind of choices serve to *explore* the action space. Alternatively the agent chooses the action such that  $a_t = \arg \max_{a_t} Q(s_t, a_t)$ , in which case the agent *exploits* its current knowledge to improve certainty on the optimal path. Learning will be the process by which the agent updates the values in  $Q$  to make better decisions in the  $1 - \epsilon$  cases where it chooses the optimal choice. Updating  $Q$  is an iterative process, and essentially follows the idea of a Bellman equation, augmented by a learning rate  $\eta \in (0, 1)$ . That is, we update  $Q$  according to:

$$Q(s_t, a_t) \leftarrow (1 - \eta)Q(s_t, a_t) + \eta(r_t + \beta \max_a Q(s_{t+1}, a)) \quad (6)$$

where  $r_t$  is the instantaneous reward  $u(a_t)$  and  $\beta$  is the discount factor [19]. To fully train the model we repeatedly expose the Q learner to the environment over  $k$  *episodes*, letting it choose actions and reap rewards while updating the  $Q$ -matrix.

Figure 1 plots the learned value map when solving the model given in (4). In period 0, no information is learned for  $s_0 < 100$  as this initial condition is fixed – thus, we only show values for  $s_0 = 100$ . The highest value is achieved from choosing  $a_0 \approx 40$ . The need to include the time step as a state becomes apparent when looking at panel C. Because the continuation value is 0 for all choices in this period, the reward is independent of the remaining stock.



**Fig 1.** Per-time period value maps learned by the Q-learner in 50,000 episodes with  $\bar{s} = 100$ ,  $T = 3$ ,  $\eta = 0.6$  and  $\beta = 0.9$ . We take  $\epsilon = 0.5$  to excessively explore the state-action space. Note in panel A we only show values for  $s_0 = 100$ .

In figure 2, we show the number of times the algorithm visits each field in  $Q$  during the 50,000 episodes that we used to generate figure 1, highlighting the usefulness of relying on nonconverged versions of  $Q$  for guiding the exploration of the state-action space.

Unlike in classical dynamic programming algorithms, the Q-learner requires fixing hyperparameters to achieve good convergence to the true value map. In particular, the number of learning episodes, the